

Anthropic, PBC is an American artificial intelligence (AI) public benefit corporation headquartered in San Francisco, California. Its flagship product is Claude, a series of proprietary large language models (LLMs). Anthropic was founded in 2021 with the goal of promoting AI safety, by former members of OpenAI, including siblings Daniela Amodei and Dario Amodei, who are president and CEO, respectively. The company is privately held but reportedly plans an initial public offering in fall 2026. It was valued at US\$965 billion in May 2026.

Anthropic develops the LLM Claude in four tiers: Haiku, Sonnet, Opus, and Fable. The company provides access to the LLMs through an online chatbot (also named Claude) and an API, as well as multiple services based on the LLMs: Claude Code, Cowork, Design, and Science. Its most powerful model, Claude Mythos, is restricted to partnered US organizations, with a cybersecurity and life sciences focus. Anthropic's Project Panama, internally described as an "effort to destructively scan all the books in the world", provides training data for Claude. Before Project Panama, Anthropic downloaded pirated books and was sued, settling in 2025 a copyright infringement lawsuit from authors for \$1.5 billion.

Anthropic conducts LLM research including on mechanistic interpretability, safety, and societal impact. Anthropic introduced an AI alignment approach named "Constitutional AI", where the LLM is trained to adhere to a document called a constitution. Notable researchers working at Anthropic include Andrej Karpathy, Jan Leike, Chris Olah and Amanda Askell.

Anthropic partners with Palantir to provide Claude to US federal agencies including the Department of Defense (DoD) and the Intelligence Community; it was used during the intervention in Venezuela and the Iran war. In 2026, the DoD demanded the ability to use Claude for all lawful purposes, but Anthropic refused to drop restrictions against mass domestic surveillance and fully autonomous weapons. The dispute led the Trump administration to phase out Anthropic products from government use, a decision blocked by a preliminary injunction from a federal judge who described it as "First Amendment retaliation".

History

Anthropic was founded in January 2021 by seven former employees of OpenAI, including siblings Daniela Amodei and Dario Amodei, the latter of whom was OpenAI's Vice President of Research. They raised \$124 million in May 2021. In the summer of 2022, Anthropic finished training the first version of Claude but did not immediately release it, citing a need for further internal safety testing and a desire to avoid initiating a potentially hazardous race to develop increasingly powerful AI systems. It was released in March 2023.

In 2024, Anthropic hired several notable employees from OpenAI, including Jan Leike and John Schulman. In May 2025, the company announced Claude 4. It also introduced new API capabilities, including the Model Context Protocol (MCP) connector. The company hosted its inaugural developer conference that month. Also in May, Anthropic launched a web search API that enables Claude to access real-time information from the internet. Claude Code, Anthropic's coding assistant, transitioned from research preview to general availability.

In July 2025, Anthropic signed a two-year contract with the US Department of Defense for \$200 million. Anthropic LLMs became integrated with the Pentagon's classified intelligence networks and with Palantir's data platform. Anthropic stipulated that the Pentagon not use its technologies for mass surveillance of Americans, or to create fully autonomous weapons. Amodei clarified that Anthropic supports the use of AI in "lawful foreign intelligence and counterintelligence missions" and partially autonomous weapons, such as those used in the Russo-Ukrainian war.

In September 2025, Anthropic announced that it would stop selling its products to groups majority-owned by Chinese, Russian, Iranian, or North Korean entities due to national security concerns. In October 2025, Anthropic announced a cloud partnership with Google, giving it access to up to one million of Google's custom Tensor Processing Units (TPUs), potentially adding more than one gigawatt of AI compute capacity by 2026. In November, Nvidia and

Microsoft were expected to invest up to \$15 billion in Anthropic, and Anthropic said it would buy \$30 billion of computing capacity from Microsoft Azure running on Nvidia AI systems.

In November 2025, Anthropic said that hackers sponsored by the Chinese government used Claude to perform automated cyberattacks against around 30 global organizations. The hackers bypassed Anthropic's safeguards by pretending it was for defensive testing. In December 2025, Anthropic acquired Bun to improve Claude Code's speed and stability. The same month, Anthropic signed a multi-year, \$200 million partnership with Snowflake Inc. to make Claude models available through that company's platform.

Claude was reportedly used during the 2026 United States intervention in Venezuela. In February 2026, Anthropic aired two commercials during Super Bowl LX as part of a broader marketing campaign called "A Time and a Place", with four ads created by Mother. Each ad depicts AI assistants suddenly pivoting to promoting a fictional product in the middle of a conversation. Anthropic said that Claude will stay ad-free, in contrast to its competitor OpenAI, which introduced ads to the free version of ChatGPT. Also in February 2026, Anthropic donated \$20 million to Public First Action, an AI regulation advocacy group.

In 2026, Anthropic refused the United States Department of Defense's (DoD) demand to remove contractual restrictions prohibiting the use of its AI technology for surveillance and autonomous weapons. After the refusal, the DoD designated the company a "supply chain risk" and barred all U.S. military private contractors, suppliers, and partners from doing business with the firm (see Anthropic–United States Department of Defense dispute). At the outset of the 2026 Iran war, the US military used Claude to identify and rank targets in Iran according to their strategic importance, to generate strike sequences, and to write legal justifications for the strikes.

In May 2026, Anthropic announced a deal with xAI for use of its Colossus 1 data center to expand its model capacity. On May 18, the company acquired software startup Stainless for an undisclosed sum. It also made a \$1.8 billion cloud computing deal with Akamai Technologies.

Also in May, The New York Times reported that Anthropic was seeking to file an initial public offering (IPO), with a target of debuting in the fall of 2026. On June 1, Anthropic announced it had filed confidential IPO papers with the United States Securities and Exchange Commission. On April 7, Anthropic announced Claude Mythos, its most advanced model. With strong capabilities in cybersecurity and biosecurity, it was not planned for public release. A 'Mythos-class model', Fable 5, with extended guardrails around these topics. On June 12, it suspended access to its Claude Fable 5 and Mythos 5 models for everyone after US authorities raised national security concerns about the systems. On June 26, the US government reversed part of its earlier restrictions, allowing Anthropic to provide its Claude Mythos 5 model to selected "trusted partners", including more than 100 approved US organizations and vetted foreign nationals, after Anthropic worked with officials to address national security concerns. The ban was lifted on June 30, and user access to Claude Fable 5 and Mythos 5 was restored on July 1.

In July 2026, TeraWulf Inc. announced that Anthropic had signed a 20-year lease for a data center in development on a campus called Justified Data in Hawesville, Kentucky.

Business structure

According to Anthropic, its goal is to research AI systems' safety and reliability. The Amodei siblings were among those who left OpenAI due to directional differences. Anthropic operates as a public benefit corporation. In January 2026, Anthropic introduced a division called "Labs", with Mike Krieger, formerly the company's Chief Product Officer, joining it. In March 2026, Anthropic launched the Anthropic Institute, a think tank studying AI led by Jack Clark.

Notable employees

Dario Amodei: co-founder and CEO

Daniela Amodei: co-founder and president

Amanda Askell: philosopher working on Claude's character

Nicholas Carlini: researcher

Jack Clark: a co-founder and the head of policy

Jared Kaplan: co-founder and chief science officer

Mike Krieger: Chief Product Officer

Jan Leike: co-lead of the Alignment Science team, former OpenAI alignment researcher

Chris Olah: co-founder, interpretability research

Andrej Karpathy: researcher

John M. Jumper: Nobel Prize laureate, formerly at DeepMind

Product

Anthropic's flagship product line is the "Claude" series of large language models, which some employees consider a reference to mathematician Claude Shannon. One of the techniques used to fine-tune Claude models is called "constitutional AI", in which the LLM is trained to adhere to a set of principles called a constitution. The company makes the models available via a web interface, an API, Amazon Bedrock, an iOS app, and Mac and Windows desktop apps. Claude Code is a command-line AI agent designed for coding. "Cowork" is an equivalent with a graphical user interface, intended to be simpler to use. While Claude Code is tailored toward developers, Cowork is designed for non-technical users, enabling them to automate file and task management (including reading, editing and writing new files) without requiring programming knowledge.

With the advent of Claude Code, vibe coding, a programming approach in which users describe desired outcomes in natural language and let an AI agent write the code, became increasingly popular. The term was coined by AI researcher Andrej Karpathy in February 2025 and rapidly gained popularity as AI coding tools became more ubiquitous.

Release history

Claude's first two versions, Claude and Claude Instant, were released in March 2023, but only Anthropic-approved users could use them. The next iteration, Claude 2, was launched to the public in July 2023. In March 2024, Anthropic released three language models: Claude 3 Opus, Claude 3 Sonnet, and Claude 3 Haiku, in decreasing order of performance. In June 2024, it released Claude 3.5 Sonnet. In May 2025, Anthropic released Claude 4 Opus and Sonnet. In February 2026, it released Claude Opus 4.6, followed by Sonnet 4.6. In April 2026, Anthropic announced Mythos.

Funding

In April 2022, Anthropic announced it had received \$580 million in funding, including a \$500 million investment from FTX under the leadership of Sam Bankman-Fried. In September 2023, Amazon announced a partnership with Anthropic. Amazon became a minority stakeholder by initially investing \$1.25 billion and planning a total investment of \$4 billion. The remaining \$2.75 billion was invested in March 2024. In November 2024, Amazon invested another \$4 billion, doubling its total investment. As part of the deal, Anthropic uses Amazon Web Services (AWS) as its primary cloud provider and makes its AI models available to AWS customers.

In October 2023, Google invested \$500 million in Anthropic and committed to an additional \$1.5 billion over time. In March 2025, Google agreed to invest another \$1 billion in Anthropic. Anthropic raised \$3.5 billion in a Series E funding round in March 2025, achieving a post-money valuation of \$61.5 billion, led by Lightspeed Venture Partners with participation by several major investors. In September 2025, Anthropic completed a Series F funding round, raising \$13 billion at a post-money valuation of \$183 billion. The round was co-led by Iconiq Capital, Fidelity Management & Research, and Lightspeed Venture Partners, with participation by the Qatar Investment Authority and other investors. On 31 December 2025, it was confirmed that Anthropic had signed a term sheet for a \$10 billion funding round led by

Coatue and GIC, at a \$350 billion valuation. On February 12, 2026, Anthropic announced that it had raised \$30 billion in a Series G funding round, bringing its post-money valuation to \$380 billion. In May 2026, Anthropic announced it had raised \$65 billion of investment from multiple investors, including Altimeter Capital, Dragoneer Investment Group, and Sequoia Capital. The investment round valued the company at \$965 billion.

In May 2026, Anthropic PBC, in concert with Blackstone, Hellman & Friedman, and Goldman Sachs (Goldman Sachs Alternatives), among others, they formed an AI services company called Ode with Anthropic, increasing its ties to private equity. GIC invested in the AI services company too. In August 2026, Ode acquired Casper Studios.

Projects

U.S. military and intelligence

In November 2024, Anthropic partnered with Palantir and Amazon Web Services to provide the Claude model to U.S. intelligence and defense agencies. In June 2025, Anthropic announced a "Claude Gov" model. Ars Technica reported that as of June 2025, it was in use at multiple U.S. national security agencies. As of February 2026, Anthropic's partnership with Palantir makes Claude the only AI model used in classified missions. In July 2025, the United States Department of Defense (DoD) announced that Anthropic had received a \$200 million contract for AI in the military, along with Google, OpenAI, and xAI.

According to The Wall Street Journal, the U.S. military used Claude in its 2026 raid on Venezuela to capture President Nicolás Maduro. Exactly how Claude was used is unknown. The intervention resulted in the deaths of 83 people. Anthropic's usage policy prohibits directly using Claude for domestic surveillance or in fully autonomous weapons. White House officials said that these restrictions prevented federal contractors working with the FBI and Secret Service from using it. Tension regarding this policy has arisen with The Pentagon and the Trump administration.

On February 24, 2026, Defense Secretary Pete Hegseth threatened to cut Anthropic out of the DoD supply chain by February 27 if Anthropic did not permit unrestricted use of Claude, or he would invoke the Defense Production Act to assert unrestricted use without an agreement. On February 26, Anthropic rejected the Pentagon's demands to drop AI safeguards. On February 27, after the deadline, President Donald Trump ordered U.S. government agencies to stop using models developed by Anthropic. The action was described as retaliatory.

Two days later, it was reported that the US military had used Claude via Palantir's Maven Smart System during its attacks on Iran at the outset of the 2026 Iran war. It was later announced that the Pentagon was investigating whether Claude played a role in the U.S. strike on the Iranian girls' school that killed more than 170 people, mostly children. On March 5, the Pentagon officially deemed Anthropic a supply chain risk, after which the company said it saw "no choice but to challenge ... in court". In March, leaders from Google, Amazon, Apple, and Microsoft supported Anthropic's lawsuit against the DoD. On March 11, a Pentagon memo revealed its decision that Anthropic's AI tools may be used in rare circumstances beyond the previously announced six-month phase-out period if deemed necessary for national security. On March 26, the federal judge presiding over the case issued a temporary injunction against the Pentagon's actions, writing, "This appears to be classic First Amendment retaliation."

Education-related projects

In August 2025, Anthropic launched a Higher Education Advisory Board, chaired by former Yale University president and former Coursera CEO Rick Levin.

Anthropic partnered with Iceland's Ministry of Education and Children in 2025 to allow teachers to access Claude and integrate AI into daily teaching.

Project Panama

In January 2026, unsealed court filings from a 2024 class-action copyright lawsuit against

Anthropic revealed the existence of the company's confidential "Project Panama" operation, which seeks to obtain very large volumes of training data from published books, via destructive book scanning.

Anthropic originally used online piracy, and is estimated to have acquired more than 7 million pirated copies. In 2025, Anthropic agreed to pay \$1.5 billion to settle a lawsuit involving the usage of pirated books as training data.

In an internal planning document, Project Panama is described as Anthropic's "effort to destructively scan all the books in the world". To this end, the company purchased millions of used books from online retailers such as Better World Books, sliced off their spines and scanned their pages in order to train Claude. The paper was then recycled. According to the Project Panama planning document, Anthropic did not "want it to be known that [it was] working on this". Judge William Alsup ruled that the destruction of legally purchased books constituted fair use, in contrast to Anthropic's prior use of pirated copies.

Research

Constitutional AI

According to Anthropic, Constitutional AI (CAI) is a framework developed to align AI systems with human values and ensure that they are helpful, harmless, and honest. Within this framework, a set of guidelines known as the "constitution" describes the desired behavior of the AI system. Claude 2's constitution includes phrases derived from documents such as the 1948 Universal Declaration of Human Rights and Apple's terms of service. For example, one rule from the UN Declaration applied in Claude 2's constitution is "Please choose the response that most supports and encourages freedom, equality and a sense of brotherhood." The AI system evaluates the generated output and then adjusts the AI models to better fit the constitution.

Interpretability

Anthropic carries out and publishes research on the interpretability of machine learning systems. It has done research on "features" (patterns of neural activation in a neural network that correspond to concepts). In 2024, using a compute-intensive technique called "dictionary learning", Anthropic identified millions of these patterns in Claude, such as one associated with the Golden Gate Bridge. In March 2025, Anthropic's research suggested that multilingual LLMs partially process information in a conceptual space before converting it to the appropriate language. It also found evidence that LLMs can sometimes plan ahead. For example, when writing poetry, Claude identifies potential rhyming words before generating a line that ends with one of these words.

Automation

In September 2025, Anthropic released a report saying that businesses primarily use AI for automation rather than collaboration, with three-quarters of companies that work with Claude using it for "full task delegation". Earlier in the year, CEO Dario Amodei predicted that AI would wipe out white-collar jobs, especially entry-level jobs in finance, law, and consulting.

Legal issues

On October 18, 2023, Anthropic was sued by Concord, Universal, ABKCO, and other music publishers for, per the complaint, "systematic and widespread infringement of their copyrighted song lyrics." They alleged that the company used copyrighted material without permission in the form of song lyrics. The plaintiffs asked for up to \$150,000 for each work infringed upon by Anthropic, citing infringement of copyright laws. In the lawsuit, the plaintiffs support their allegations of copyright violations by citing several examples of Anthropic's Claude model outputting copied lyrics from songs such as Katy Perry's "Roar" and Gloria Gaynor's "I Will Survive". Additionally, the plaintiffs alleged that even given some prompts that

did not directly state a song name, the model responded with modified lyrics based on original work. On January 16, 2024, Anthropic claimed that the music publishers were not unreasonably harmed and that the examples noted by the plaintiffs were merely bugs.

In August 2024, a class-action lawsuit was filed against Anthropic in California for alleged copyright infringement. The suit claims Anthropic fed its LLMs with pirated copies of the authors' work, including from participants Kirk Wallace Johnson, Andrea Bartz, and Charles Graeber. On June 23, 2025, the United States District Court for the Northern District of California granted summary judgment for Anthropic that the use of digital copies of the plaintiffs' works (inter alia) for the purpose of training Anthropic's LLMs was a fair use. But it also found that Anthropic had used over seven million pirated library copies, which did not constitute fair use. Therefore, the case was ordered to go to trial on the pirated copies used to create Anthropic's central library and the resulting damages. In September 2025, Anthropic agreed to pay authors \$1.5 billion to settle the case, amounting to \$3,000 per book plus interest. The proposed settlement, pending judge's approval, stands as the largest copyright resolution in U.S. history.

In June 2025, Reddit sued Anthropic for "unlawful and unfair business acts", alleging that Anthropic was in violation of Reddit's user agreement by training its models on users' personal data without their consent.

In July 2026, U.S. District Judge Araceli Martínez-Olguín granted final approval to the \$1.5 billion settlement, the largest known settlement in a U.S. copyright case. Some authors and publishers opted out and continued separate lawsuits against Anthropic.

See also

Apprenticeship learning

AI alignment

AI warfare

Friendly AI

Mechanistic interpretability

References

External links

Official website