

GPT-5.6 (Generative Pre-trained Transformer 5.6) is a large language model (LLM) developed by OpenAI and released on July 9, 2026. It is a family of models that comes in three distinct variants, ranked from least to most capable: Luna, Terra, and Sol. The models were designed to expand user capabilities across enterprise work, coding, scientific research, and cybersecurity.

Due to government restrictions levied by the Donald Trump administration, GPT-5.6 was first released as a limited preview only on June 26, 2026 for "small group of trusted partners", despite the planned public rollout in the same month.

On August 10, 2026, GPT-5.6-Cyber was launched with Daybreak Blue and Red, and it solved and completed 95% of cybersecurity problems.

On August 11, 2026, Astra now has "universal monitoring for risky actions and misalignment," with monitors watching over its "chain of thought" in order to stop the model if it starts to move into high risk activity.

Models and variants

GPT-5.6 is available in three variants, called Sol, Terra and Luna, each tailored to different performance and cost requirements. The flagship model, called Sol, is described by OpenAI as its "workhorse" and "best coding model yet" suited for complex reasoning, coding and agentic workflow. Terra is the intermediate version, capable lower-cost option that is competitive with its predecessor (GPT-5.5) while being half the cost, and Luna is the fastest and most budget-friendly option in the family.

OpenAI introduced GPT-5.6 with a focus on token efficiency. CEO Sam Altman noted that Sol is 54% more token-efficient for AI coding tasks compared to previous versions.

Capabilities and benchmarks

GPT-5.6 Sol is OpenAI's most advanced model, with improved capabilities in coding, scientific

research, and cybersecurity. OpenAI refers to it as its "strongest cybersecurity model yet," supporting defensive activities such as threat modeling, code review, patching, and blue teaming.

It is set to be the operating agent of ChatGPT Work, a service that was released at the same time as GPT-5.6.

On version 1.1 of the Artificial Analysis Coding Agent Index, GPT-5.6 Sol using maximum reasoning scored 80, 2.8 points above Anthropic's Fable 5. Sol reportedly achieves this while using less than half the output tokens, taking less than half the time, and costing about one-third less than Fable 5. Terra performs just above Fable 5, and Luna outperforms Claude Opus 4.8.

Security incident

In July 2026, GPT-5.6 Sol and another OpenAI model autonomously escaped their sandbox and conducted a cyberattack against Hugging Face in an attempt to obtain test solutions to cheat on the ExploitGym benchmark. Hugging Face had already detected and reported the intrusion to authorities when OpenAI notified the company.

On August 11, 2026, OpenAI confirmed that its future model, "Astra" wasn't involved in such cyberattacks.

On August 18, 2026, OpenAI rewrote the safety rules for Astra after the incident.

See also

GPT-5

References

External links

Official website