

LTX is a family of open source artificial intelligence video foundation models developed by Lightricks, and first released in November 2024. The latest generation, LTX-2, includes releases such as LTX-2.3 and LTX-2.5 and generates synchronized video and audio from text, image, and video inputs. They were preceded by LTX Video, which was released in 2024 as the company's first text-to-video model.

LTX-2 is part of the LTX family of video generation models, which form the core technology, alongside LTX Studio, of the LTX ecosystem.

History

Origins: LTX Video (2024–2025)

In November 2024 Lightricks publicly released its first text-to-video model, LTX Video. It was a 2-billion parameter model, available as open source.

In May 2025 Lightricks launched LTXV-13b, a version with 13-billion parameters. Two months later, the model broke the 60 second barrier for generated video.

Release of LTX-2 (2025)

In October 2025 Lightricks announced its latest model, and renamed it LTX-2. The model was described as capable of generating synchronized audio and video at native 4K resolution and up to 50 frames per second (fps), using a variety of conditions and prompts, including text-to-video and image-to-video.

Google highlighted the fact that LTX-2 was trained on its infrastructure, and saying it was "The first open source AI video generation model, powered by Google Cloud".

Upon its release it was ranked in the top-3 models for image-to-video creation by Artificial Analysis, behind Kling 3.5 by Kling AI and Veo 3.1 by Google. Its text-to-image option was ranked 7th.

In addition to its open-source release, Lightricks offers API access to LTX-2, allowing

developers to generate videos from text and image prompts through a hosted service without running the model locally.

Open Source Release (2026)

In January 2026, Lightricks officially released the full open-source version of LTX-2, making the model's complete codebase, weights, and associated tooling publicly available.

In March 2026 the company released LTX-2.3, which was accompanied by a desktop video editor enabling the entire model to run locally on consumer hardware.

In August 2026 LTX-2.5 was released, highlighted by its move towards world models used for robotics training, and the ability to generate 10 second videos in under 7 seconds. LTX-2.5 reported to add native multishot generation and a pretrained checkpoint intended for domain-specific fine-tuning. The release continued LTX's development as a world model, with applications including robotics and physical AI.

Technical features

Advancements over LTX Video

LTX-2 builds upon the LTX Video architecture with several major improvements:

Unified audio-video generation producing synchronized dialogue, ambient, and motion

Native 4K rendering

50-fps output for cinematic motion

Three operational modes (Fast, Pro, Ultra)

More efficient diffusion pipelines enabling high fidelity on consumer GPUs

Core capabilities

Text-to-video generation

Image-to-video generation

Video-to-video generation and video conditioning

Multimodal audiovisual synthesis

High-resolution spatial and temporal coherence

Configurable quality/performance settings

Open-source distribution of weights and datasets

Native multishot generation

HDR and OpenEXR input and output workflows

Diffusion-based video decoding

Custom Gemma 4 12B text encoder

Pretrained LTX-2.5 checkpoint for domain-specific fine-tuning

Reception

Initial reception to LTX-2 was broadly positive, with several technology and media outlets highlighting its open-source approach and multimodal capabilities. Open Source For You described LTX-2 as “one of the first AI video systems to combine 4K output, synchronized audio, and an open model release,” noting that it positioned Lightricks as a significant competitor to proprietary systems such as OpenAI's Sora and Google's Veo.

IEA Green said that the model “could rewrite the AI filmmaking game,” emphasizing that its 50-fps rendering and unified audio-video generation made it suitable for professional studios and independent creators alike.

AI News characterized LTX-2 as a “major step forward in the democratization of cinematic-quality video generation,” praising its consumer-grade hardware efficiency and multi-tier generation modes, while also noting ongoing challenges in long-form temporal stability.

FinancialContent reported strong interest among creative agencies, attributing the attention

to Lightricks' decision to release model weights and datasets, which reviewers said enabled "a level of transparency not typically seen in commercial AI video models."

Benchmarks and rankings

Upon release, LTX-2 ranked third for image-to-video creation in the Artificial Analysis benchmark, behind Kling 3.5 and Veo 3.1, while its text-to-video option ranked seventh. As of early 2026, it was the highest-ranked open-source model in the benchmark.

Limitations

Some early reviewers also pointed out quality limitations. The Ray3 technical review noted occasional inconsistencies in lip-sync and motion tracking during long scenes, though it stated these were "in line with the challenges faced by all current AI video diffusion models" and expected to improve with continued iteration.

Like other diffusion-based video generators, LTX-2 can produce artifacts in complex multi-person scenes and may struggle with precise text rendering within generated video.

See also

Sora (text-to-video model)

Veo (text-to-video model)

Generative artificial intelligence

Kling (text-to-video model)

Runway (company)

Diffusion model

World model (artificial intelligence)

References

External links

Official website

[1]

[2]

Official Github Repo