

Qwen (also known as Tongyi Qianwen, Chinese: 通义千问; pinyin: Tōngyì Qiānwèn) is a family of large language models developed by Alibaba Cloud.

Models

Qwen

Alibaba launched a beta of Qwen in April 2023 under the name Tongyi Qianwen, then opened it for public use in September 2023 after regulatory clearance.

The model's architecture was based on the Llama architecture developed by Meta AI. In December 2023, it released its 72B and 1.8B models for download, while Qwen 7B weights were released in August.

Qwen2

Qwen2 was released in June 2024, and in September it released some of its models with open weights, while keeping its most advanced models proprietary. Qwen2 contains both dense and sparse models.

In November 2024, QwQ-32B-Preview, a model focusing on reasoning similar to OpenAI's o1, was released under the Apache 2.0 License, although only the weights were released, not the dataset or training method. QwQ has a 32K token context length and performs better than o1 on some benchmarks. It was also in November 2024 that the Accio application was launched.

The Qwen-VL series is a line of visual language models that combines a vision transformer with an LLM. Alibaba released Qwen2-VL with variants of 2 billion and 7 billion parameters.

In January 2025, Qwen2.5-VL was released with variants of 3, 7, 32, and 72 billion parameters. All models except the 72B variant are licensed under the Apache 2.0 license. Qwen-VL-Max is Alibaba's flagship vision model as of 2024, and is sold by Alibaba Cloud at a cost of US\$0.41 per million input tokens.

Alibaba has released several other model types such as Qwen-Audio and Qwen2-Math. In

total, it has released more than 100 open weight models, with its models having been downloaded more than 40 million times. Fine-tuned versions of Qwen have been developed by enthusiasts, such as "Liberated Qwen", developed by San Francisco-based Abacus AI, which is a version that responds to any user request without content restrictions.

On 29 January 2025, Alibaba launched Qwen2.5-Max.

On 24 March 2025, Alibaba launched Qwen2.5-VL-32B-Instruct as a successor to the Qwen2.5-VL model. It was released under the Apache 2.0 license.

On 26 March 2025, Qwen2.5-Omni-7B was released under the Apache 2.0 license and made available through chat.qwen.ai, as well as platforms like Hugging Face, GitHub, and ModelScope. The Qwen2.5-Omni model accepts text, images, videos, and audio as input and can generate both text and audio as output, allowing it to be used for real-time voice chatting.

Qwen3

On 28 April 2025, the Qwen3 model family was released, with all models licensed under the Apache 2.0 license. The Qwen3 model family includes both dense and MoE models. The sizes of the dense models include 0.6B, 1.7B, 4B, 8B, 14B, and 32B, and the MoE models include 30B-A3B (30B with 3B activated parameters) and 235B-A22B (235B with 22B activated parameters). They were trained on 36 trillion tokens in 119 languages and dialects.

The Qwen3 collection includes:

Qwen3, which supports switching between thinking and non-thinking mode

Qwen3 Base, the pretrained base model

Qwen3 Instruct, which supports only non-thinking mode

Qwen3 Thinking, which supports only thinking mode

In addition to open-weights models, Qwen3 also includes Qwen3-Max, a proprietary large model with over 1T parameters, trained with about 36T tokens, that is available through API,

and Qwen3-Max-Thinking, its reasoning variant that can generate text, pictures, or video.

In February 2026, Alibaba released the open-weights Qwen3.5 and the proprietary Qwen3.5-Plus. Alibaba stated that Qwen 3.5 is able to operate desktop and mobile applications. Qwen3.5-Omni and Qwen3.6-Plus were released in April 2026 as proprietary; access to these tools is limited to the chatbots' websites and the Alibaba cloud platform. The Qwen3.6 model was released under the Apache License in the same month. Alibaba released the proprietary Qwen3.7 models in Max and Plus variants in May and June 2026, respectively. Alibaba previewed its 2.4-trillion-parameter model Qwen3.8-Max in July 2026 and announced that it would release the weights of the model. The announcement occurred a few days after Moonshot AI released its competing Kimi K3 model. The cloud version of Qwen3.8-Max was released on 3 August 2026. It uses a sparse mixture-of-experts architecture, with approximately 95 billion parameters active per forward pass, and supports a context window of up to one million tokens. Alibaba released the weights of Qwen3.8-Max as Qwen3.8-2.4T-A95B on 12 August 2026, and the open-weights model omits certain features of the cloud model, such as image input and a non-thinking mode. Qwen3.8-2.4T-A95B's license requires model providers generating more than US\$50 million in revenue within 12 months to obtain a commercial license with Alibaba.

On 14 August 2026, Alibaba followed up on the release of Qwen3.8-Max with the release of Qwen3.8-72B, which includes both the image input and non-thinking mode omitted in the larger model release. In contrast with the more restrictive lines of Qwen3.8-Max, this model was released with the Apache 2.0 license.

List of models

History

In a January 2026 update of the Qwen mobile application, the Alibaba Group connected the chatbot to the company's ecosystem, starting with food-service delivery. The company

announced plans to allow users to assign ecosystem-related-tasks to more company platforms such as Taobao and Fliggy and to help consumers with errands, including phone calls and document processing.

Former Qwen AI model division head, Lin Junyang, resigned in March 2026 after the release of Qwen3.5 and Qwen3.5-Plus, becoming the latest of three executives to have left Alibaba this year. Amid concern that this could mean a shift away from research and open-source artificial intelligence, Alibaba said it will continue its focus on open source.

Later the same month, a company announcement revealed the formation of a new AI business unit, Alibaba Token Hub. The new unit will supervise AI-related work, led by Alibaba CEO Eddie Wu and includes Tongyi Large Model Business Unit, MaaS Business Line, Qwen, Wukong, and AI Innovation. The Token Hub is further supported by senior leaders Zhou Jingren, as chief AI architect, and CTO Wu Zeming. Led by Jingren, Alibaba's Tongyi Large Model Business Unit, formerly known as Tongyi Laboratory, was restructured to specialize in the development of Qwen AI models.

Distribution

Many Qwen models are distributed under the free and open-source Apache License, the source-available Qwen License, or the non-commercial Qwen Research License.

On the other hand, proprietary Qwen models are served through Alibaba Cloud.

Open-source strategy

AI Singapore announced in November 2025 that they will replace LLaMA models with the open-sourced Qwen framework as the base of the Sea-Lion LLM model

Reception

There are over 200,000 variations of Qwen's open-source AI models on the Hugging Face's

model list; Qwen3-VL-2B-Instruct has surpassed 18 million downloads globally. The platform contains many small, specialized Qwen derivatives by third-party developers. Models can also be found on Alibaba's ModelScope, which was created for greater open-source accessibility due to China's restriction of Hugging Face in 2022. By May 2026, the Qwen app had 234 million users.

The U.S.-China Economic and Security Review Commission reported that the Chinese approach in developing open-source AI models, such as Qwen, has been critical to China's ability to overcome constraints in compute and to expand its capabilities in real-world data curation. This enhanced capacity in data curation is perceived to be important to China's ability to integrate AI into research, manufacturing, and robotics sectors.

Related products

Accio

Accio is an AI native application that is built upon Qwen and is used to generate market insights and answer sourcing questions for Alibaba's business to business e-commerce site. The tool is able to automate labor-intensive tasks like data collection and trend tracking.

Accio Work

Accio Work is an agentic AI platform that was launched by Alibaba International in March 2026 to fully automate e-commerce and sourcing tasks for small and medium businesses.

It uses multiple specialized AI agents that can conduct market research, track trends, find suppliers, compare quotes, generate listings, and manage basic store operations, reducing manual data collection and monitoring work.

The system is described as built on Alibaba's trade data and supports several large language models, including Qwen for different agent roles.

Because it does not rely only on Qwen as a large language model, it would be false to assume

that it wholly depends on Qwen. It is often described as running on a hybrid NLP pipeline where Qwen handles semantic understanding and natural-language reasoning.

See also

List of large language models

Lists of open-source artificial intelligence software

References

External links

Official website

Qwen on ModelScope

Qwen on Hugging Face

Qwen on GitHub