

Text watermarking is a technique for embedding hidden information within textual content to verify its authenticity, origin, or ownership. With the rise of generative AI systems using large language models (LLM), there has been significant development focused on watermarking AI-generated text. Potential applications include detecting fake news and academic cheating, and excluding AI-generated material from LLM training data. With LLMs the focus is on linguistic approaches that involve selecting words to form patterns within the text that can later be identified. The results of the first reported large-scale public deployment, a trial using Google's Gemini chatbot, appeared in October 2024: users across 20 million responses found watermarked and unwatermarked text to be of equal quality. Research on text watermarking began in 1997.

See also

Digital watermarking

References

== Further reading ==