

Velvet AI is a family of multilingual generative artificial intelligence models developed by Almage, an Italian AI company. The Velvet models, including Velvet 14B and Velvet 2B, are foundational large language models (LLMs) designed and developed in Italy on Almage's proprietary architecture. They were trained on the Leonardo supercomputer managed by CINECA and have been released in open-weight format. The Velvet models are designed to be energy-efficient. They support multiple languages, with an emphasis on Italian.

Models

Velvet 2B

Velvet 2B, the model with 2 billion parameters, was trained in Italian and English on over 2 trillion tokens of data.

Velvet 14B

Velvet 14B was trained on over 4 trillion tokens across six languages, with Italian comprising approximately 23% of the data. In addition to linguistic data, Velvet "incorporates over 400 billion tokens from more than 100 programming languages to facilitate more structured inferences."

The development of Velvet AI reflects Almage's strategic investment in creating high-performance, energy-efficient AI solutions that align with European regulatory frameworks. The models are ready to be used on major market platforms in the cloud, on-premise, and on the edge, and are integrated into Almage's AIWave platform.

Velvet 25B

Velvet 25B is a large language model (LLM) in Almage's Velvet family. In general natural language processing terminology, LLMs are pretrained models learned by predicting tokens from context and used for conditional text generation. Almage announced Velvet 25B on 14

October 2025, together with the multimodal text-to-speech model Velvet Speech 2B. The company described Velvet 25B as a 25-billion-parameter model optimized for text processing in all 24 official languages of the European Union and designed for advanced reasoning in complex environments.

According to Almagest, Velvet 25B was trained natively across European languages rather than by using English as a reference language. The company states that training began from more than 15 trillion tokens and ended with more than 7 trillion tokens, and that the model has a 128,000-token context window, a 148,000-token vocabulary, six million instruction examples used for specialization, and a June 2025 data cutoff. Almagest said that the long-context architecture was intended to preserve consistency and accuracy across distant passages in large documents, including legal texts, scientific dossiers and legislative acts. On its technology page, Almagest presents Velvet 25B as suitable for large and complex texts in regulated or specialized domains and states that it can be deployed on a single GPU.

The model is described by Almagest as supporting instruction following, information extraction, multi-step dynamic reasoning, structured output generation, function calling, machine translation, textual entailment, classification, question answering, multi-turn conversation, summarization, paraphrasing and retrieval-augmented generation (RAG). RAG is a method for combining parametric generation with retrieved non-parametric information sources. Almagest also said that Velvet 25B can be customized through reinforcement learning and integrated with agentic systems and external knowledge sources. A later company blog post listed healthcare, European law and public administration, education and culture, manufacturing and industry, and customer care among the domains represented in the training and customization pipeline.

Almagest positioned Velvet 25B as an efficient enterprise model intended to reduce energy consumption and operating costs. The company said that the model could operate in cloud, on-premises and edge environments, and reported the use of optimization methods including

adaptive computation time, post-training quantization, knowledge distillation and 4D parallelization, with claimed GPU-cluster efficiency above 90%. The model is integrated into Almagest's AIWave platform, which includes no-code and low-code tools for conversational agents, knowledge navigation and case management.

Coverage by *la Repubblica* described Velvet 25B as part of an Italian, non-Anglocentric approach to generative AI, emphasizing privacy-preserving algorithms and compliance with European regulation. The Italian industry publication *AI4Business* listed Almagest among Italy's main enterprise AI developers and noted that Velvet 25B was presented for the 24 official EU languages with emphasis on reasoning, efficiency and operation on comparatively limited hardware, while cautioning that the performance claims came from the company. Mark Up similarly presented Velvet 25B as a European multilingual AI model designed to run on limited infrastructure and to support enterprise use cases, including retail applications.

References

See also

CINECA

Istituto Italiano per l'Intelligenza Artificiale (AI4I)